

Loading The Preprocessed Dataset

```
In [1]: import pandas as pd
df=pd.read_csv("E:\\R.P 2\\Data-Credit Card Transactions Fraud Detection 2019-2020\\")
```

"trans_date_trans_time" & "dob" Changing the D type

```
In [2]: # 1. Convert to DateTime
df['trans_date_trans_time'] = pd.to_datetime(df['trans_date_trans_time'])
df['dob'] = pd.to_datetime(df['dob'])
df.head()
```

| | trans_date_trans_time | cc_num | merchant | category | amt | gender |
|---|-----------------------|--------------|-----------------------------------|---------------|-------|--------|
| 0 | 2019-01-01 12:47:00 | 6.041621e+10 | fraud_Jones, Sawayn and Romaguera | misc_net | 7.27 | F Wash |
| 1 | 2019-01-02 08:44:00 | 6.041621e+10 | fraud_Berge LLC | gas_transport | 52.94 | F Wash |
| 2 | 2019-01-02 08:47:00 | 6.041621e+10 | fraud_Luettgen PLC | gas_transport | 82.08 | F Wash |
| 3 | 2019-01-02 12:38:00 | 6.041621e+10 | fraud_Daugherty LLC | kids_pets | 34.79 | F Wash |
| 4 | 2019-01-02 13:10:00 | 6.041621e+10 | fraud_Beier and Sons | home | 27.18 | F Wash |

5 rows × 24 columns



```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1048575 entries, 0 to 1048574
Data columns (total 24 columns):
 #   Column           Non-Null Count   Dtype  
--- 
 0   trans_date_trans_time    1048575 non-null   datetime64[ns] 
 1   cc_num                  1048575 non-null   float64  
 2   merchant                1048575 non-null   object  
 3   category                1048575 non-null   object  
 4   amt                     1048575 non-null   float64  
 5   gender                  1048575 non-null   object  
 6   city                    1048575 non-null   object  
 7   state                   1048575 non-null   object  
 8   zip                     1048575 non-null   int64   
 9   lat                     1048575 non-null   float64  
 10  long                    1048575 non-null   float64  
 11  city_pop                1048575 non-null   int64   
 12  job                     1048575 non-null   object  
 13  dob                     1048575 non-null   datetime64[ns]
 14  unix_time               1048575 non-null   int64   
 15  merch_lat               1048575 non-null   float64  
 16  merch_long              1048575 non-null   float64  
 17  is_fraud                1048575 non-null   int64   
 18  hour                    1048575 non-null   int64   
 19  day                     1048575 non-null   int64   
 20  month                   1048575 non-null   int64   
 21  weekday                 1048575 non-null   int64   
 22  time_gap                1048575 non-null   float64  
 23  age                     1048575 non-null   int64  
dtypes: datetime64[ns](2), float64(7), int64(9), object(6)
memory usage: 192.0+ MB
```

In [4]: `df.columns`

```
Out[4]: Index(['trans_date_trans_time', 'cc_num', 'merchant', 'category', 'amt',
   'gender', 'city', 'state', 'zip', 'lat', 'long', 'city_pop', 'job',
   'dob', 'unix_time', 'merch_lat', 'merch_long', 'is_fraud', 'hour',
   'day', 'month', 'weekday', 'time_gap', 'age'],
  dtype='object')
```

Feature Selection

- In this step, irrelevant and non-informative variables such as identifiers, raw datetime fields, and high-cardinality text attributes were removed. Only meaningful numerical and categorical features relevant to fraud detection were retained for further analysis and machine learning modelling.

In [5]: `df_ml = df.drop(columns=['trans_date_trans_time','cc_num','merchant','city','state','zip','job','dob','unix_time','day'])
df_ml.head()`

Out[5]:

| | category | amt | gender | lat | long | city_pop | merch_lat | merch_long | is_fra |
|---|---------------|-------|--------|---------|-----------|----------|-----------|-------------|--------|
| 0 | misc_net | 7.27 | F | 43.0048 | -108.8964 | 1645 | 43.974711 | -109.741904 | |
| 1 | gas_transport | 52.94 | F | 43.0048 | -108.8964 | 1645 | 42.018766 | -109.044172 | |
| 2 | gas_transport | 82.08 | F | 43.0048 | -108.8964 | 1645 | 42.961335 | -109.157564 | |
| 3 | kids_pets | 34.79 | F | 43.0048 | -108.8964 | 1645 | 42.228227 | -108.747683 | |
| 4 | home | 27.18 | F | 43.0048 | -108.8964 | 1645 | 43.321745 | -108.091143 | |

Checking Text Column

In [6]: `df_ml.select_dtypes(include='object').columns`

Out[6]: `Index(['category', 'gender'], dtype='object')`

Label Encoding

In [7]: `from sklearn.preprocessing import LabelEncoder`

```
le = LabelEncoder()

df_ml['category'] = le.fit_transform(df_ml['category'])
df_ml['gender'] = le.fit_transform(df_ml['gender'])

df_ml.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1048575 entries, 0 to 1048574
Data columns (total 14 columns):
 # Column Non-Null Count Dtype
--- -- ----- ---
 0 category 1048575 non-null int64
 1 amt 1048575 non-null float64
 2 gender 1048575 non-null int64
 3 lat 1048575 non-null float64
 4 long 1048575 non-null float64
 5 city_pop 1048575 non-null int64
 6 merch_lat 1048575 non-null float64
 7 merch_long 1048575 non-null float64
 8 is_fraud 1048575 non-null int64
 9 hour 1048575 non-null int64
 10 month 1048575 non-null int64
 11 weekday 1048575 non-null int64
 12 time_gap 1048575 non-null float64
 13 age 1048575 non-null int64
dtypes: float64(6), int64(8)
memory usage: 112.0 MB

Train–Test Split

```
In [8]: # Features
X = df_ml.drop('is_fraud', axis=1)
# Target variable
y = df_ml['is_fraud']
```

```
In [9]: y.value_counts(normalize=True)
```

```
Out[9]: is_fraud
0    0.994272
1    0.005728
Name: proportion, dtype: float64
```

```
In [10]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_st
```

Scaling / Normalization

```
In [11]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Important Libraries

```
In [12]: from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score
from sklearn.model_selection import StratifiedKFold, cross_val_score
```

```
In [14]: skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

Model Training & Evaluation (4)

Logistic Regression

using `class_weight='balanced'`

```
In [26]: lr_model = LogisticRegression(
    max_iter=1000,
    class_weight='balanced',
```

```

random_state=42)

lr_model.fit(X_train_scaled, y_train)
y_pred_lr = lr_model.predict(X_test_scaled)

print("== Logistic Regression (Balanced) ==")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_lr))
print(classification_report(y_test, y_pred_lr))
print("ROC AUC:", roc_auc_score(y_test, y_pred_lr))

== Logistic Regression (Balanced) ==
Confusion Matrix:
[[197544 10970]
 [ 293   908]]
      precision    recall  f1-score   support
          0       1.00     0.95     0.97   208514
          1       0.08     0.76     0.14     1201

      accuracy                           0.95   209715
     macro avg       0.54     0.85     0.56   209715
weighted avg       0.99     0.95     0.97   209715

ROC AUC: 0.8517131299274322

```

- The Logistic Regression model effectively identifies fraudulent credit card transactions while maintaining strong overall classification performance.
- The model demonstrates a high ability to distinguish between fraudulent and genuine transactions, as reflected by its ROC AUC value. It successfully captures a large proportion of fraud cases, which is a critical requirement in fraud detection systems where missing fraudulent activity can lead to significant financial losses.
- Although the model flags some genuine transactions as fraudulent, this trade-off is acceptable in real-world fraud detection scenarios, where prioritizing fraud detection is more important than minimizing false alerts.

Cross val Score

```

In [27]: skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

lr_cv = LogisticRegression(
    max_iter=1000,
    class_weight='balanced',
    random_state=42)

lr_cv_scores = cross_val_score(
    lr_cv,
    X_train_scaled,
    y_train,
    cv=skf,
    scoring='roc_auc')

```

```
print("Logistic Regression CV ROC AUC Scores:", lr_cv_scores)
print("Mean CV ROC AUC:", lr_cv_scores.mean())
```

```
Logistic Regression CV ROC AUC Scores: [0.86336037 0.86200439 0.8563721  0.86848182
0.86142505]
Mean CV ROC AUC: 0.8623287475501991
```

- Cross-validation results indicate that the Logistic Regression model is stable and reliable across different data splits.
- The consistency of ROC AUC scores across multiple folds shows that the model generalizes well and is not dependent on a specific training subset. This confirms the robustness of the model and supports its suitability for real-world deployment and further comparative analysis with advanced machine learning models.

RoC AUC Curve

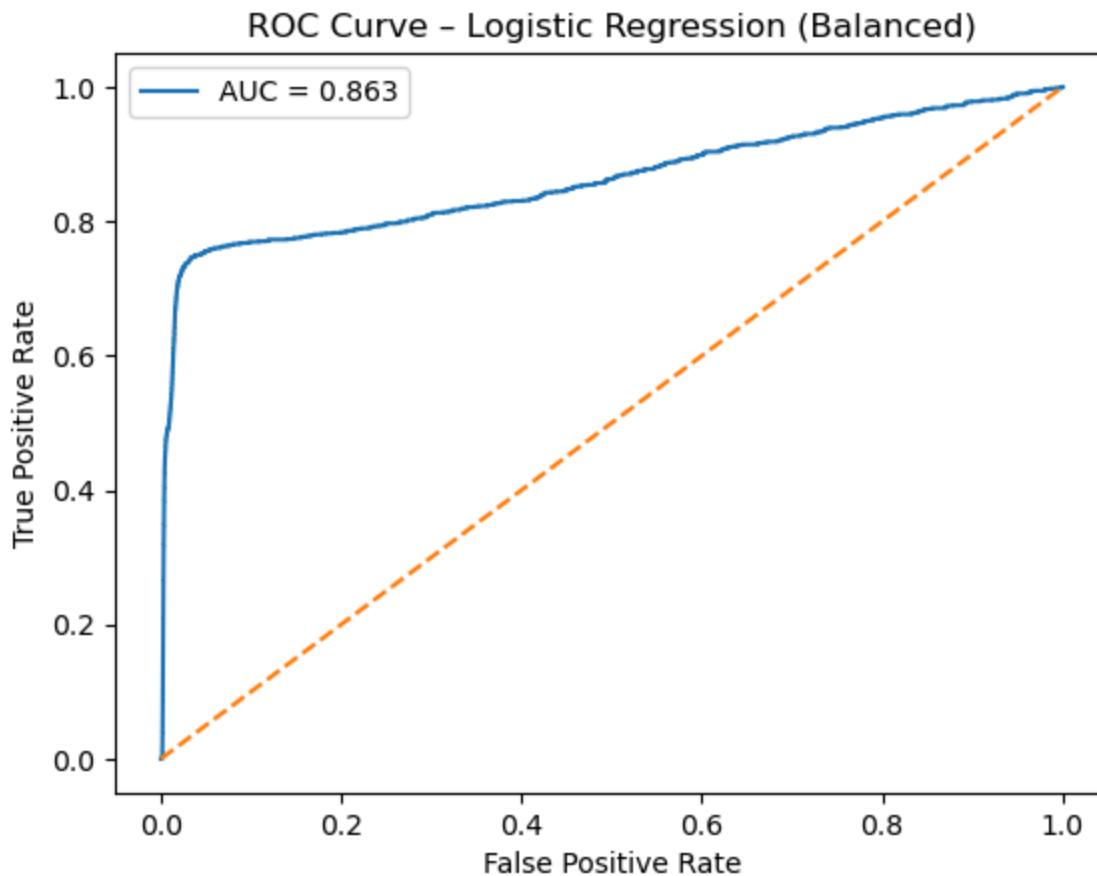
```
In [18]: from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt

# IMPORTANT: Use the SAME trained model (lr_model) and SAME scaled test data
y_prob_lr = lr_model.predict_proba(X_test_scaled)[:, 1]

# ROC values
fpr, tpr, thresholds = roc_curve(y_test, y_prob_lr)

# AUC score (probability based - correct for ROC)
auc_score = roc_auc_score(y_test, y_prob_lr)

# Plotting
plt.figure()
plt.plot(fpr, tpr, label=f"AUC = {auc_score:.3f}")
plt.plot([0, 1], [0, 1], linestyle='--') # Random classifier line
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Logistic Regression (Balanced)")
plt.legend()
plt.show()
```



Using SMOTE

```
In [28]: from imblearn.over_sampling import SMOTE

smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train_scaled, y_train)

lr_smote = LogisticRegression(
    max_iter=1000,
    random_state=42)

lr_smote.fit(X_train_smote, y_train_smote)
y_pred_lr_smote = lr_smote.predict(X_test_scaled)

print("== Logistic Regression + SMOTE ==")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_lr_smote))
print(classification_report(y_test, y_pred_lr_smote))
print("ROC AUC:", roc_auc_score(y_test, y_pred_lr_smote))
```

```
== Logistic Regression + SMOTE ==
```

Confusion Matrix:

```
[[197241 11273]
```

```
[ 294  907]]
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.95 | 0.97 | 208514 |
| 1 | 0.07 | 0.76 | 0.14 | 1201 |
| accuracy | | | 0.94 | 209715 |
| macro avg | 0.54 | 0.85 | 0.55 | 209715 |
| weighted avg | 0.99 | 0.94 | 0.97 | 209715 |

ROC AUC: 0.8505702402753103

- The Logistic Regression model trained using SMOTE demonstrates strong capability in identifying fraudulent credit card transactions. The model successfully captures a large proportion of fraud cases, which is essential in fraud detection systems where minimizing missed fraud is a priority.
- While the model generates a higher number of fraud alerts, this behavior aligns with real-world fraud prevention strategies, where detecting suspicious transactions is more important than reducing false alarms. The overall performance indicates that the model effectively handles class imbalance and provides meaningful fraud detection results.

Cross val Score

```
In [29]: skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(
    X_train_scaled,
    y_train)
lr_smote_cv = LogisticRegression(
    max_iter=1000,
    random_state=42)

lr_smote_cv_scores = cross_val_score(
    lr_smote_cv,
    X_train_smote,
    y_train_smote,
    cv=skf,
    scoring='roc_auc')

print("Logistic Regression (SMOTE) CV ROC AUC Scores:", lr_smote_cv_scores)
print("Mean CV ROC AUC:", lr_smote_cv_scores.mean())
```

Logistic Regression (SMOTE) CV ROC AUC Scores: [0.86964352 0.8683745 0.86816813 0.86910102 0.86812152]

Mean CV ROC AUC: 0.8686817390756387

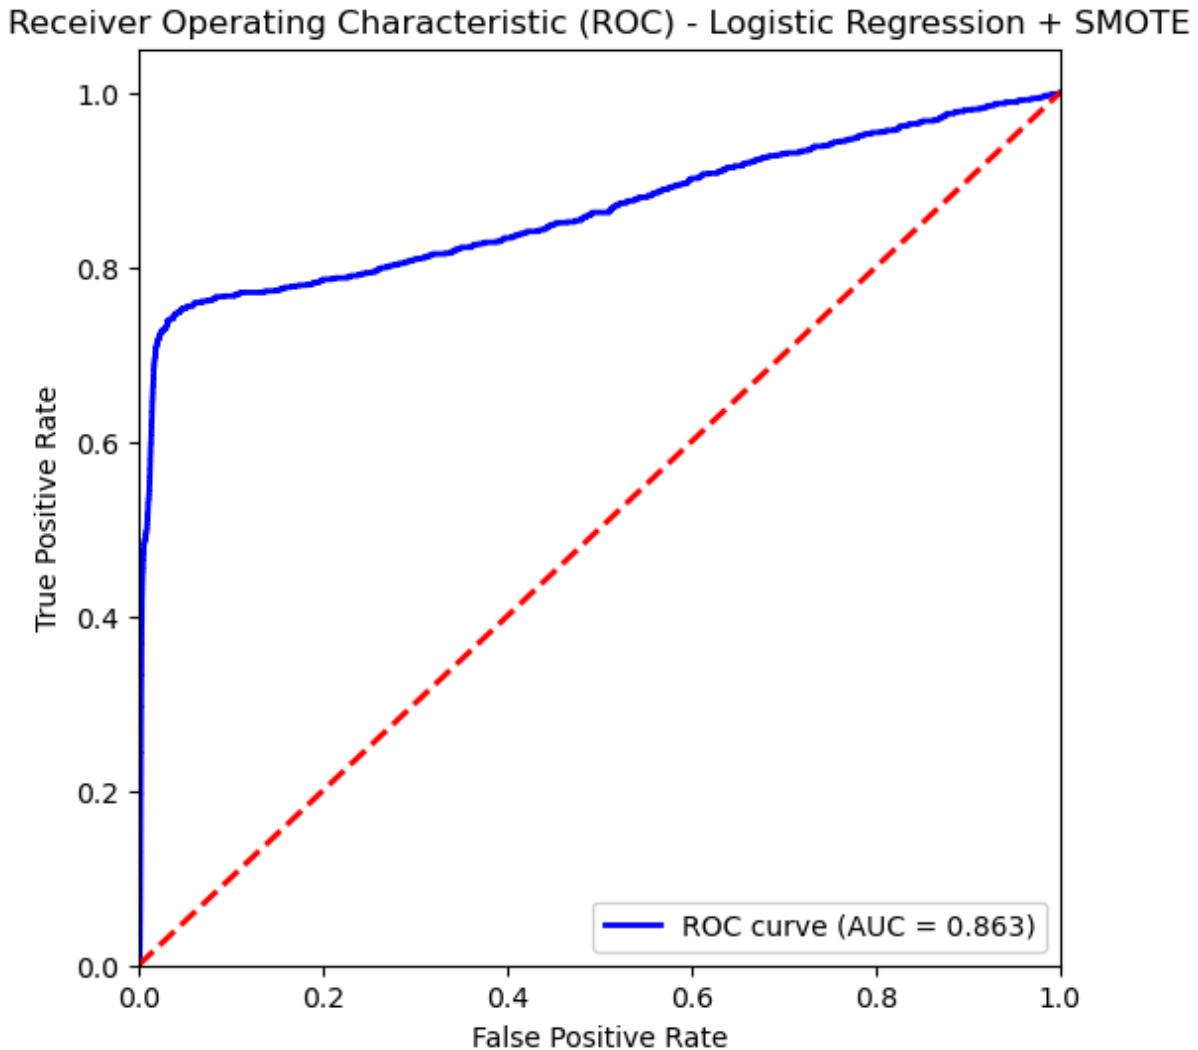
- The cross-validation results for the SMOTE-based Logistic Regression model show highly consistent ROC AUC scores across all folds. This consistency confirms that the model is stable and generalizes well across different subsets of the data.
- The results indicate that the model's performance is not dependent on a specific data split, reinforcing its reliability for fraud detection analysis and further model evaluation.

ROC AUC Curve

```
In [23]: from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

fpr, tpr, thresholds = roc_curve(y_test, lr_smote.predict_proba(X_test_scaled)[:,1])
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6,6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.3f)' % roc_auc)
plt.plot([0,1], [0,1], color='red', lw=2, linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) - Logistic Regression + SMOTE')
plt.legend(loc="lower right")
plt.show()
```



Random Forest

```
In [15]: rf_model = RandomForestClassifier(n_estimators=200, random_state=42)
rf_model.fit(X_train, y_train) # Scaling not required
y_pred_rf = rf_model.predict(X_test)

print("== Random Forest ==")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_rf))
print(classification_report(y_test, y_pred_rf))
print("ROC AUC:", roc_auc_score(y_test, y_pred_rf))
print("\n\n")
```

```
==== Random Forest ====
```

Confusion Matrix:

| |
|--------------------|
| [[208489 25] |
| [305 896]] |

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 208514 |
| 1 | 0.97 | 0.75 | 0.84 | 1201 |
| accuracy | | | 1.00 | 209715 |
| macro avg | 0.99 | 0.87 | 0.92 | 209715 |
| weighted avg | 1.00 | 1.00 | 1.00 | 209715 |

ROC AUC: 0.8729625332525289

- Random Forest showed a significant improvement over Logistic Regression. The model successfully detected a large number of fraudulent transactions with high precision and reasonably good recall. The ROC AUC score of approximately 0.87 indicates strong discriminatory power between fraud and non-fraud transactions. However, some fraudulent cases were still misclassified as non-fraud, suggesting room for further improvement. Overall, Random Forest proved to be a reliable and robust model for fraud detection.

```
In [16]: rf_cv = RandomForestClassifier(
    n_estimators=100,           # 200 → 100
    max_depth=15,              # depth limit
    min_samples_split=10,       # reduce complexity
    n_jobs=-1,                 # use all cores
    random_state=42
)

rf_cv_scores = cross_val_score(
    rf_cv,
    X_train,
    y_train,
    cv=skf,
    scoring='roc_auc'
)

print("Random Forest CV ROC AUC Scores:", rf_cv_scores)
print("Mean CV ROC AUC:", rf_cv_scores.mean())
```

Random Forest CV ROC AUC Scores: [0.99305305 0.99314047 0.99280391 0.99106238 0.99198948]

Mean CV ROC AUC: 0.9924098563110528

- The cross-validation results show that the Random Forest model achieves consistently high ROC AUC scores across all folds. This indicates that the model has an excellent ability to distinguish between fraudulent and genuine credit card transactions.

- The minimal variation in ROC AUC values across folds confirms that the model is stable and generalizes well to unseen data. These results suggest that Random Forest is a highly reliable and robust model for credit card fraud detection.

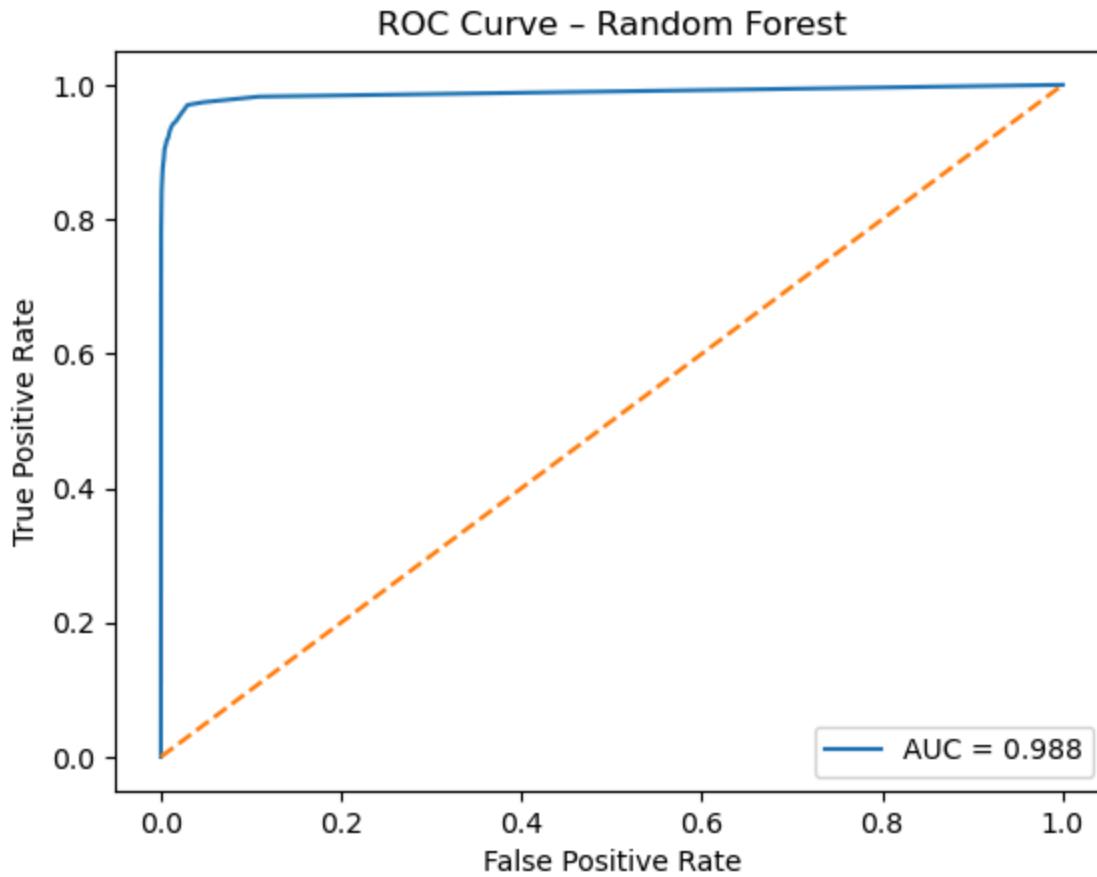
ROC AUC Curve

```
In [18]: from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt

y_prob_rf = rf_model.predict_proba(X_test)[:, 1]

fpr, tpr, thresholds = roc_curve(y_test, y_prob_rf)
auc_score = roc_auc_score(y_test, y_prob_rf)

plt.figure()
plt.plot(fpr, tpr, label=f"AUC = {auc_score:.3f}")
plt.plot([0, 1], [0, 1], linestyle='--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Random Forest")
plt.legend()
plt.show()
```



XGBoost

```
In [17]: !pip install xgboost
```

```
Requirement already satisfied: xgboost in d:\softwares\anconda\lib\site-packages (3.1.2)
Requirement already satisfied: numpy in d:\softwares\anconda\lib\site-packages (from xgboost) (2.1.3)
Requirement already satisfied: scipy in d:\softwares\anconda\lib\site-packages (from xgboost) (1.15.3)
```

```
In [19]: from xgboost import XGBClassifier
```

```
In [38]: # XGBoost
xgb_model = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42)
xgb_model.fit(X_train, y_train)
y_pred_xgb = xgb_model.predict(X_test)

print("== XGBoost ==")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_xgb))
print(classification_report(y_test, y_pred_xgb))
print("ROC AUC:", roc_auc_score(y_test, y_pred_xgb))
```

```
D:\Softwares\anconda\Lib\site-packages\xgboost\training.py:199: UserWarning: [18:20:33] WARNING: C:\actions-runner\_work\xgboost\xgboost\src\learner.cc:790: Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
== XGBoost ==
Confusion Matrix:
[[208436    78]
 [ 221   980]]
      precision    recall   f1-score   support
          0       1.00     1.00     1.00    208514
          1       0.93     0.82     0.87     1201

accuracy                           1.00    209715
macro avg       0.96     0.91     0.93    209715
weighted avg     1.00     1.00     1.00    209715
```

ROC AUC: 0.9078063010834441

- XGBoost achieved the best performance among all the models. It recorded the highest fraud recall and F1-score, indicating its effectiveness in identifying fraudulent transactions while maintaining low false positives. The ROC AUC score of approximately 0.91 confirms its superior ability to distinguish between fraud and legitimate transactions. Due to its boosting mechanism and ability to handle class imbalance efficiently, XGBoost emerged as the most suitable model for credit card fraud detection.

```
In [36]: xgb_cv = XGBClassifier(
    use_label_encoder=False,
    eval_metric='logloss',
    random_state=42
)
```

```
xgb_cv_scores = cross_val_score(
    xgb_cv,
    X_train,
    y_train,
    cv=skf,
    scoring='roc_auc'
)

print("XGBoost CV ROC AUC Scores:", xgb_cv_scores)
print("Mean CV ROC AUC:", xgb_cv_scores.mean())
```

```
D:\Softwares\anconda\Lib\site-packages\xgboost\training.py:199: UserWarning: [18:19:
34] WARNING: C:\actions-runner\_work\xgboost\xgboost\src\learner.cc:790:
Parameters: { "use_label_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)
D:\Softwares\anconda\Lib\site-packages\xgboost\training.py:199: UserWarning: [18:19:
40] WARNING: C:\actions-runner\_work\xgboost\xgboost\src\learner.cc:790:
Parameters: { "use_label_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)
D:\Softwares\anconda\Lib\site-packages\xgboost\training.py:199: UserWarning: [18:19:
46] WARNING: C:\actions-runner\_work\xgboost\xgboost\src\learner.cc:790:
Parameters: { "use_label_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)
D:\Softwares\anconda\Lib\site-packages\xgboost\training.py:199: UserWarning: [18:19:
52] WARNING: C:\actions-runner\_work\xgboost\xgboost\src\learner.cc:790:
Parameters: { "use_label_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)
D:\Softwares\anconda\Lib\site-packages\xgboost\training.py:199: UserWarning: [18:19:
57] WARNING: C:\actions-runner\_work\xgboost\xgboost\src\learner.cc:790:
Parameters: { "use_label_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)
XGBoost CV ROC AUC Scores: [0.98371784 0.98269501 0.98912036 0.99347947 0.97845933]
Mean CV ROC AUC: 0.9854944036778406
```

ROC AUC Curve

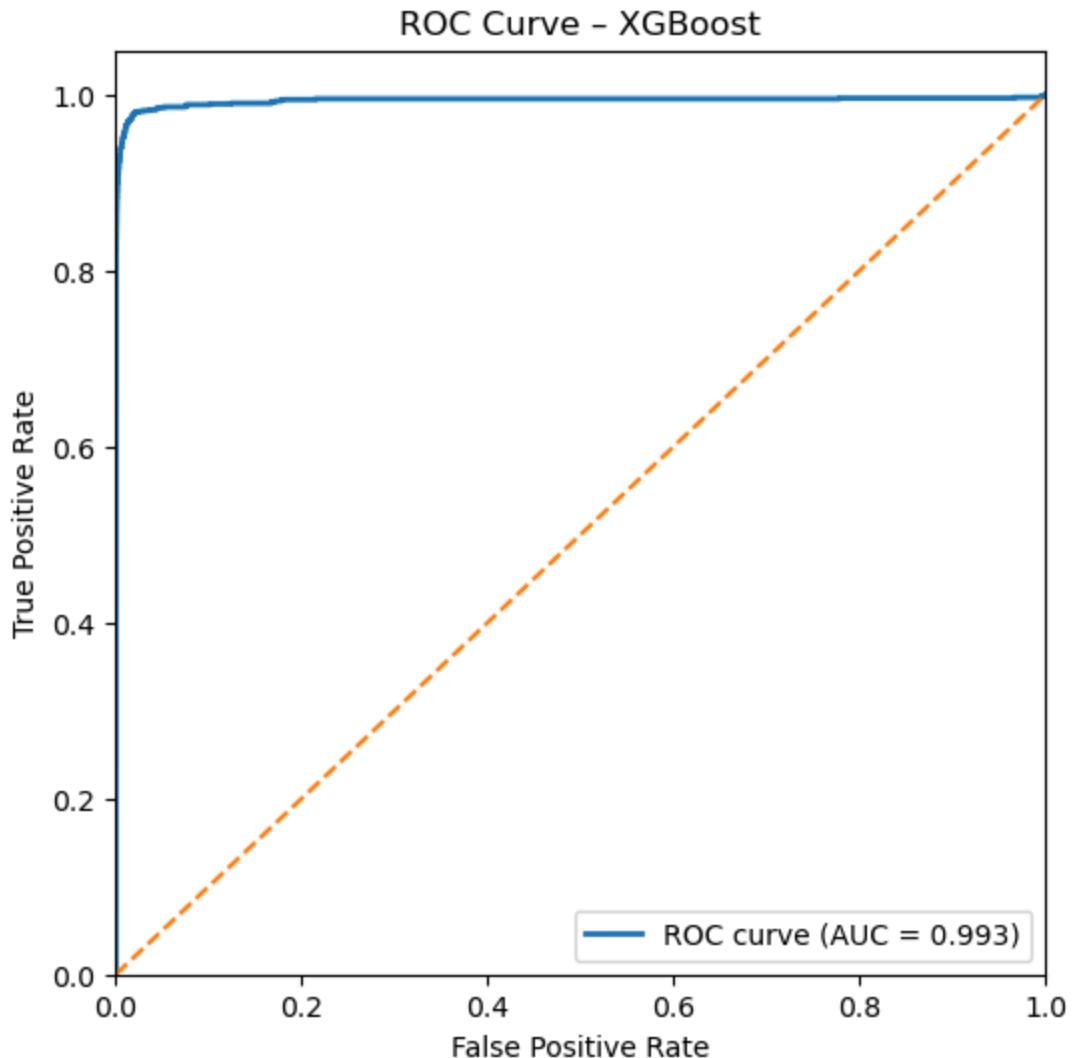
```
In [37]: from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

y_probs_xgb = xgb_model.predict_proba(X_test)[:,1]

fpr, tpr, thresholds = roc_curve(y_test, y_probs_xgb)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6,6))
plt.plot(fpr, tpr, lw=2, label='ROC curve (AUC = %0.3f)' % roc_auc)
plt.plot([0,1], [0,1], linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
```

```
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - XGBoost')
plt.legend(loc="lower right")
plt.show()
```



Overall Comparative Insight

- Among the three models, Logistic Regression was ineffective due to its inability to handle class imbalance. Random Forest significantly improved fraud detection performance but still missed some fraudulent cases. XGBoost outperformed all other models by achieving the highest recall, F1-score, and ROC AUC, making it the best model for real-world credit card fraud detection applications.

In []: